

Mean Surprise $:= -\frac{1}{|\mathcal{X}_i|} \sum_{x \in \mathcal{X}_i} \log p(x \mid \theta_{\text{RL}})$ for different base models

